

# Part 1

## COS314 Artificial Intelligence (Ch. 1–3) – Exam Summary and Key Concepts

### Executive Summary

Artificial Intelligence (AI) is the science and engineering of creating intelligent machines that can perform tasks typically requiring human intelligence. The foundations of AI trace back to pivotal questions about whether machines can **think or demonstrate intelligence**. Alan Turing's **Imitation Game (Turing Test)** was an early proposal to evaluate machine intelligence by judging if a computer's responses were indistinguishable from a human's. John Searle's **Chinese Room** argument later challenged the notion of understanding in AI, illustrating that symbol processing alone (as a computer does) might simulate comprehension without *true* understanding. These philosophical touchstones set the stage for AI's development.

AI has since evolved into various **key areas**: from **expert systems** that encode human knowledge in rules to solve specialized problems, to **natural language processing (NLP)** enabling computers to interpret and generate human language, to **artificial neural networks (ANNs)** inspired by brain architecture and capable of learning from data, and **evolutionary algorithms** that use principles of natural selection to “evolve” solutions. Classic problem domains like game playing (e.g. chess, Go), theorem proving, and robotics have driven AI research and led to remarkable milestones (such as IBM's Deep Blue defeating the chess world champion and Google DeepMind's AlphaGo mastering Go through self-play). These areas collectively demonstrate AI's range from symbolic reasoning to data-driven learning.

A unifying theme in AI problem-solving is **state space search**. Complex tasks are framed as searching through a **state space** (a set of possible states and transitions) to reach a goal state from a starting state. This framework applies to **classic puzzles and problems** like the **Traveling Salesman Problem**, **Towers of Hanoi**, and the **8-Puzzle**, which are often used to benchmark search strategies. **Uninformed search methods** (blind search) like **Depth-First Search (DFS)** and **Breadth-First Search (BFS)** systematically explore states without domain-specific guidance, whereas **informed (heuristic) search methods** like **Best-First**

**Search, Hill-Climbing, and A\*** leverage problem-specific knowledge to find solutions more efficiently. The choice of search strategy affects completeness (will it find a solution if one exists?), optimality (does it find the best solution?), and efficiency (time and memory usage).

A fundamental concept in informed search is the **heuristic function** – an educated guess of how close a given state is to a goal. Good heuristics drastically improve search efficiency. For instance, in the 8-puzzle, counting how many tiles are out of place or the total Manhattan distance of tiles from their goal positions guides algorithms like A *to the solution faster than brute force. However, heuristics must be designed carefully: an **admissible** heuristic (one that never overestimates the true cost to reach the goal) ensures optimal solutions when used with A.* Developing effective heuristics often involves insightful simplifications of the problem.

In summary, Chapters 1–3 cover AI's early philosophical questions, principal subfields, the formulation of problems as state-space searches, and the array of search techniques from basic exhaustive searches to sophisticated heuristic-guided algorithms. These concepts form the core toolkit for AI problem-solving and are likely exam topics. Mastery of definitions (e.g. what is the Turing Test?), comparisons (e.g. DFS vs BFS, or expert systems vs neural networks), and applications (e.g. using A *for route-finding, using expert systems in medical diagnosis*) is essential. Understanding theoretical underpinnings – such as why A finds optimal paths or how evolutionary algorithms mimic biology – will not only help in exam settings but also deepen your appreciation of how AI methods relate to human intelligence and natural processes.

## Historical Foundations of AI

**Birth of AI and Fundamental Questions:** The field of AI formally began in the mid-20th century (notably the 1956 Dartmouth workshop) when researchers asked questions like “What is intelligence?” and “Can machines think or learn like humans?”. Early experiments focused on making computers perform tasks such as game playing and theorem proving, which were seen as requiring “intelligence.” These efforts highlighted two key components of AI research: the **ability to search** through possible solutions and the **representation of knowledge** in a form a machine can use.

**Turing Test (Imitation Game):** In 1950, Alan Turing proposed a thought experiment to define a criterion for machine intelligence. In the **Turing Test**, a human evaluator engages in natural language conversations with two hidden counterparts – one human and one machine. If the evaluator cannot reliably tell which is which, the machine is said to have demonstrated

intelligence indistinguishable from a human. This test shifted the question from “*Can machines think?*” to “*Can machines behave in a way that’s indistinguishable from thinking humans?*”. The Turing Test has been highly influential, framing AI’s goal as producing human-like responses. However, it also emphasizes **behavioral indistinguishability** over internal understanding – a fact often discussed in exams. (*An interesting modern insight: Some AI chatbots have “passed” limited versions of the Turing Test by cleverly mimicking human conversational patterns, though critics note this doesn’t guarantee true understanding.*)

**Chinese Room Argument:** Proposed by philosopher John Searle in 1980, the Chinese Room is a rebuttal to the idea that passing the Turing Test equates to real understanding. Searle imagines a person who does not speak Chinese locked in a room with a big rulebook for Chinese symbol manipulation. Chinese speakers outside the room slide questions in Chinese under the door; the person inside uses the rulebook to look up responses and slides out answers in Chinese. To an external observer, it appears the person (or “system”) inside understands Chinese, because the responses are appropriate. In reality, the person doesn’t understand the content at all – they are simply following syntactic rules. **Searle’s point:** a computer executing a program could similarly manipulate symbols (0s and 1s, or words) to produce valid outputs without any understanding or consciousness. This thought experiment highlights the difference between *simulating* understanding and *possessing* understanding (syntax vs semantics). For exam purposes, remember that the Chinese Room challenges the notion that intelligent behavior implies actual comprehension, suggesting that passing a Turing Test might not be sufficient evidence of true intelligence.

Together, the Turing Test and Chinese Room cover the **philosophical foundations** of AI. They underscore issues of **intelligence versus imitation** and **syntax versus semantics**. In an exam setting, you might be asked to explain these concepts, compare them, or discuss what they imply about the nature and goals of AI. An important takeaway is that AI’s history is not just technical but also philosophical – these debates guide how we evaluate AI systems beyond just performance metrics.

## Core Areas of Artificial Intelligence

AI is a broad field encompassing many sub-disciplines. Below are key areas of AI, each with distinct goals and methodologies, that have emerged (many of these are likely to be referenced in exams with examples):

- **Game Playing:** Using AI to play games such as chess, checkers, or Go at a high level. Early AI research used games as a testbed (for example, Arthur Samuel’s checkers

program in the 1950s learned to play competitively). Game-playing AI must handle complex **search spaces**; classic techniques include minimax search with heuristics and pruning (e.g., alpha-beta pruning in chess). *Real-world parallel:* AI milestones like IBM's **Deep Blue** (which defeated the world chess champion in 1997 using brute-force search and domain heuristics) and **AlphaGo** (which defeated Go champions in 2016–2017 by combining deep neural networks with tree search and reinforcement learning) demonstrate how games drive AI innovation. These systems illustrate how raw computing power plus clever heuristics or learning can achieve superhuman performance in well-defined domains.

- **Theorem Proving:** Automating logical reasoning to prove mathematical theorems or solve logic puzzles. AI theorem provers (such as the early *Logic Theorist* or modern satisfiability solvers) use formal logic representations and search through the space of possible proofs. This area emphasizes **knowledge representation** (formalizing facts/axioms) and **search algorithms** to navigate inference rules. Though not as famous as game AIs, theorem proving has contributed to fields like software verification and automated reasoning.
- **Expert Systems:** These are AI programs that emulate the decision-making ability of a **human expert** in a specific domain. They were among the most successful AI applications in the 1970s-1980s. An expert system consists of a **knowledge base** (a set of domain-specific facts and rules, often encoded as IF-THEN rules) and an **inference engine** that applies the rules to known facts to deduce new information or make decisions. For example, *MYCIN* was an expert system for medical diagnosis of blood infections, using dozens of rules to recommend treatments as a human doctor would. *Key characteristics:* Expert systems perform well in **narrow domains** (they are only as good as the expert knowledge encoded), and they typically do not learn from data – they rely on human experts to update their rules. On exams, an **expert system** might be defined, and you might be asked about its components or to contrast it with learning-based approaches (for instance, unlike a neural network which learns from examples, an expert system's knowledge is manually curated). **Insight:** In modern AI, expert systems have in part been superseded by machine learning, but their legacy survives in how we design systems that incorporate explicit rules (e.g., fraud detection systems combining rules with learned models). They highlight the **symbolic AI** approach – explicitly encoding human knowledge – as opposed to the subsymbolic approach of neural networks.
- **Natural Language Processing (NLP) / Understanding:** NLP is the area focused on enabling computers to **understand, interpret, and generate human language**. This includes tasks like machine translation, speech recognition, chatbots, sentiment analysis, and more. NLP is challenging because human language is complex and ambiguous. Early NLP systems used rule-based parsing and grammars (e.g., SHRDLU in the 1970s could manipulate virtual blocks via language commands), while modern NLP heavily employs **statistical and neural network models** (for example, the sequence-to-sequence models for translation or transformers for language understanding). A practical example is

voice assistants (Siri, Alexa) that convert speech to text, interpret the query, and formulate a spoken answer. For exams, remember that NLP involves both **syntax** (structure of language) and **semantics** (meaning). An interesting recent development is the rise of **large language models (LLMs)** like GPT, which can generate human-like text; these models demonstrate how learning from vast datasets can yield language understanding far beyond what early rule-based NLP achieved. *(Unconventional insight: Modern NLP models draw inspiration from cognitive linguistics and even neuroscience – e.g., transformer architectures have attention mechanisms somewhat analogous to human focus of attention in processing information.)*

- **Modelling Human Performance:** This area attempts to create systems that simulate or support human cognitive processes and skills. This can include cognitive modeling (building computer models that perform tasks in ways similar to how humans do) or applications like **intelligent tutoring systems (ITS)** that mimic human tutors. For instance, an ITS might model a student's knowledge state and adapt its teaching strategy accordingly, much like a human tutor. These systems often integrate knowledge representation, inference, and sometimes natural language dialogue to interact with users. In essence, this area blurs the line between AI and cognitive science, emphasizing understanding *how* humans solve problems so the AI can mirror or assist that process.
- **Planning and Robotics:** AI planning involves generating sequences of actions that achieve specified goals, often in physical or simulated environments. Early planning algorithms (like STRIPS in 1971) could plan actions for simple robots or problem domains (e.g., a robot navigating rooms to fetch an object). Robotics combines perception (using sensors to understand the world), planning (deciding on actions), and execution (movement/control) – all of which are AI problems on their own. An example of planning in practice is path planning for a mobile robot or route planning for navigation systems. Modern robotics uses AI for tasks such as localization (estimating position), mapping environments (SLAM algorithms), and even dynamic motion planning for bipeds (like Honda's ASIMO robot). In an exam, you might not dive deep into robotics algorithms, but be aware that **planning** is a key AI capability allowing autonomous decision-making over time.
- **Artificial Neural Networks (ANNs):** ANNs are computing models inspired by the brain's neural structure. An ANN consists of many simple connected processing units (artificial "neurons") that adjust connection strengths (weights) to learn patterns from data. Neural networks are the foundation of **machine learning** approaches that have gained prominence since the 1990s (and especially in the last decade with **deep learning**). They excel at tasks like image and speech recognition, where writing explicit rules is infeasible. Key concepts include learning algorithms (like backpropagation for training multi-layer networks) and the idea that knowledge is distributed across many weighted connections rather than encoded as discrete rules. **Analogy (Neuroscience Insight):** Just as networks of biological neurons in the brain strengthen or weaken synapses through experience, ANNs adjust weights through training on examples, gradually "learning" the task. Modern **deep neural**

**networks** (with many layers) have achieved human-level or better performance in vision (e.g., classifying images, self-driving car vision) and language tasks, which were historically very challenging for symbolic AI. In contrast to expert systems' explicitly coded knowledge, neural nets *learn* implicit internal representations. An exam question might ask for the basic idea of a neural network, or to compare it with symbolic AI; highlighting this difference and the inspiration from biology would be important points. *(Recent development insight: The success of deep learning has led to AI systems that, for example, can diagnose diseases from medical images or translate speech in real time – capabilities that earlier AI approaches struggled with. This underscores how ANNs have become central in AI today.)*

- **Evolutionary Algorithms and Nature-Inspired Computing:** This category of AI techniques is inspired by biological processes. **Evolutionary algorithms (EAs)**, such as Genetic Algorithms, simulate natural selection by evolving a population of candidate solutions. Each candidate (individual) has a set of parameters (analogous to chromosomes) that can mutate or recombine with others. Over successive generations, the population “evolves” better solutions according to a fitness function (which evaluates how good a solution is). For example, an EA could evolve an optimal schedule, design a neural network architecture, or solve an engineering optimization problem without explicit human design – instead, trial and error and selection do the work. **Analogy (Evolutionary Biology Insight):** Think of trying to solve a complex problem as akin to evolving a creature adapted to an environment – many random variations are tested, and the fittest are kept to breed the next generation. Over time, good traits accumulate. EAs are particularly useful for problems where the search space is rugged and poorly understood (so traditional methods might get stuck easily). They can sometimes find creative or unconventional solutions (famously, evolutionary algorithms have even been used to design antenna shapes for NASA spacecraft and came up with designs that were unusual but highly effective!). Other nature-inspired methods include **ant colony optimization** (inspired by how ants find shortest paths via pheromones) and **swarm intelligence** (inspired by flocking behavior, used in routing and robotics). In exams, focus on the concept that evolutionary algorithms perform search via *random variation and selection*, and note how this contrasts with systematic, deterministic search like A\* or BFS. They illustrate a different approach to AI – one that **stochastically searches** for good solutions and is robust to complex search landscapes.
- **Agent-Based Technology:** In AI, an *agent* is an entity that perceives its environment and takes actions to achieve goals. Agent-based systems can be simple (a single autonomous agent like a thermostat adjusting temperature) or complex (multi-agent systems where multiple AI agents interact or cooperate/compete, such as a simulated market with trading agents, or teams of robotic soccer players). Key topics include agent architectures (reflex agents, goal-based agents, learning agents) and communication in multi-agent systems. An example is a swarm of drones coordinating to map a disaster area – each drone is an intelligent agent making local decisions, but together they achieve a larger objective. Agents

unify many AI concepts (perception, decision-making, learning) under one framework. While an exam might not delve deeply into agent theory in these chapters, recognizing that modern AI often deploys *agents* (software or robots) that operate autonomously is useful.

- **Uncertainty Reasoning:** Many real-world AI problems involve uncertainty – incomplete or noisy information. Uncertainty reasoning deals with models and algorithms that can handle ambiguity, probabilistic outcomes, and incomplete knowledge. This includes **probabilistic reasoning** (like Bayesian networks or Markov decision processes) and **fuzzy logic** (reasoning with degrees of truth rather than binary true/false). For instance, an AI medical diagnosis system might use probability to weigh different possible illnesses given symptoms, since symptoms can be caused by multiple diseases. Understanding basic ideas of probability in AI (like updating beliefs with Bayes' theorem, or the concept of an agent maximizing expected utility) is often expected by upper-level AI courses, though details might be beyond chapters 1–3. In summary, this area emphasizes that AI must often **reason under uncertainty** rather than in a perfectly known environment.
- **Generative AI:** A newly prominent area (mentioned in passing in the notes) where AI systems learn to generate new data samples similar to the training data. Examples include generative adversarial networks (GANs) creating realistic images, and large language models generating human-like text. This area is a fusion of neural networks and probabilistic modeling. Although generative AI might be outside the classic scope of an intro AI exam, it's a **real-world development** showing how far AI has advanced – AIs can now create art, write articles, or compose music by learning patterns from existing data. It's a modern testament to AI's capabilities, built on foundations from the other areas (for example, combining neural networks with optimization techniques). For exam prep, it's enough to be aware that generative models are an application of neural networks and learning, reflecting AI's growing ability not just to recognize patterns but also to create.

**Note:** These areas often overlap. For instance, AlphaGo's success in *game playing* combined neural networks (for evaluating positions, a product of deep learning) with search (Monte Carlo Tree Search planning moves) and reinforcement learning (trial-and-error learning of strategies). Likewise, a **self-driving car** is an agent that uses vision (neural nets for image recognition), planning (to navigate streets), and decision-making under uncertainty (probabilistic reasoning about what other cars will do). When studying, try to recall examples of AI systems and identify which areas and techniques they draw upon – this holistic understanding can yield *unconventional insights* and impress exam graders.

## Classic AI Problems

Throughout AI, certain problem scenarios are considered “classic” because they encapsulate common challenges and serve as standard tests for AI algorithms. You should be familiar with the following classic problems, their setup, and why they are interesting in AI:

- **Traveling Salesman Problem (TSP):** *Problem statement:* A salesman needs to visit a set of cities exactly once and return to the starting city. Given the pairwise distances (or travel costs) between cities, the goal is to find the shortest possible round-trip route. TSP is famous as a combinatorial optimization problem – the number of possible tours grows factorially with the number of cities (for  $N$  cities, there are  $(N-1)!/2$  possible tours, which becomes astronomically large even for moderate  $N$ ). TSP is **NP-hard**, meaning no efficient algorithm is known to find the optimal tour for all cases, and it's believed that none exists (this is related to the P vs NP question). In AI, TSP exemplifies a state space where brute-force search is impractical beyond small sizes, so heuristics, approximations, or clever search (like branch-and-bound or genetic algorithms) are applied. **Real-world parallels:** variations of TSP appear in logistics (route planning for deliveries), circuit board drilling, and DNA sequencing. For exam purposes, understand that TSP requires searching through permutations of cities – a huge search space – and thus is used to illustrate the need for heuristic or approximate methods. You might be asked how an AI would approach TSP (e.g., using a heuristic like the nearest neighbor for a greedy tour, or using branch and bound to prune bad routes).
- **Towers of Hanoi:** *Problem statement:* We have three pegs (rods) and a set of disks of different sizes stacked on one peg in decreasing size order (largest at bottom, smallest on top). The task is to move all disks from one peg to another, one disk at a time, never placing a larger disk on top of a smaller disk. The classic puzzle uses 3 pegs and (often) 3 disks for demonstration, but the legendary version mentions 64 golden disks in a Tibetan monastery (with the myth that when the monks finish transferring all 64 disks, the world will end!). The **state space** here is defined by the distribution of disks on pegs. For even a small number of disks, the number of states is large (with 3 pegs and  $n$  disks, there are  $3^n$  possible states). The optimal solution for  $n$  disks requires  $2^n - 1$  moves, which grows exponentially. Towers of Hanoi is important in AI as a benchmark for **recursive problem solving and search depth**. Many search algorithms can be applied to it (it's a puzzle with a clear goal state). In an exam, you might be asked to outline the state representation (e.g., which peg each disk is on), or to explain why a brute-force search would be inefficient. It's also a good example of how a **search strategy** like DFS would systematically find a solution (in fact, DFS naturally yields the standard recursive solution for Towers of Hanoi). Fun fact/insight: Towers of Hanoi solutions relate to binary numbers and Gray codes – showing a connection between this puzzle and mathematical patterns, which is sometimes an interesting aside in AI classes.
- **8-Puzzle (Sliding Tile Puzzle):** *Problem description:* We have a  $3 \times 3$  grid with 8 numbered tiles and one empty space. A tile adjacent to the empty space can slide into it, swapping the empty location. Given an initial configuration of tiles and a desired goal configuration (often the tiles in numerical order), the task is to rearrange the tiles from the start to goal by sliding tiles. The 8-puzzle's state space consists of all reachable tile configurations (there are  $9! = 362,880$  possible arrangements of 9 spaces, but only half are reachable



from a given configuration due to parity constraints, so about 181,440 reachable states per starting position). This is a quintessential **search problem**: each move changes the state, and we seek a shortest path to the goal. **Why it matters in AI:** The 8-puzzle is small enough to experiment with, but large enough that naive search will struggle, which motivated the development of heuristics. In fact, *A search was famously applied to the 15-puzzle (a larger 4x4 version) in the 1960s using heuristics like Manhattan distance (total distance of tiles from their goal positions) to guide the search. The 8-puzzle is commonly used to illustrate how different search strategies perform: BFS can find the shortest solution but will examine many states, DFS might get lost without finding the solution (or find a very long suboptimal sequence), whereas A with a good heuristic finds the solution efficiently.* For an exam, you should be able to describe how to represent the 8-puzzle state (e.g. a 3x3 array of tile positions, with the blank as a special symbol) and possibly give an example of a heuristic (like the number of misplaced tiles or Manhattan distance). **Insight:** The 8-puzzle also highlights the concept of **state space complexity** – there’s a huge number of possibilities, yet the structure of the puzzle (sliding tiles) allows intelligent algorithms to prune the search. It exemplifies why AI search algorithms are needed: a simple brute-force would have to explore up to hundreds of thousands of states, but A\* can solve typical instances by exploring only a fraction of that by focusing on states that look “closer” to the goal.

These classic problems serve as **standard testbeds** for AI search strategies. They each have a clear initial state and goal state, and a set of legal moves (operators). Importantly, they demonstrate different challenges:

- TSP emphasizes **optimization** and the combinatorial explosion of possibilities.
- Towers of Hanoi emphasizes **depth of search** and structured problem spaces.
- 8-puzzle emphasizes **branching factor** and the need for heuristics in large search spaces.

Knowing these examples helps ground abstract concepts. For instance, if asked about heuristics, you might cite the 8-puzzle’s Manhattan distance heuristic as a concrete example. Or if comparing search methods, you could discuss how BFS or DFS would navigate the Tower of Hanoi or 8-puzzle state space differently. Always tie back the theoretical concept to one of these concrete problems if possible – it shows you understand how the theory applies in practice.

## State Space Representation (Graphs vs Trees)

Before we can apply any search algorithm, we need to **formulate the problem as a state space**. A **state space** is an abstract representation of all possible states (configurations) that the problem can be in, along with the transitions (actions or moves) that lead from one state to another. Formally, a state space can be described by:

- A **initial state** (or set of possible initial states).
- One or more **goal states** (or a goal condition that defines success).
- A set of **operators or actions** that map from one state to another (legal moves).
- Optionally, a **cost** for each move (for problems where some solutions are “cheaper” or better than others).

We can think of the state space as a directed graph where:

- **Nodes** represent states.
- **Edges** represent actions leading to new states.

Depending on the problem, this graph can sometimes be conceptualized as a **tree** (especially if each state is typically reached by a unique path, or if we intentionally avoid revisiting states). More generally, it's a graph because there may be multiple ways to reach the same state from the start, or loops that can return to a prior state.

**Graph vs Tree Search:** If each state were unique to a single path (no revisiting), the search space would form a tree. In practice, many problems have **redundant paths** or **cycles**. For example, in the 8-puzzle, you could move a tile left then right, returning to the original state – a cycle. As a **tree**, this would be represented as distinct nodes, even though it's the same configuration revisited. Treating the state space as a graph acknowledges that states can repeat. **Why this matters:** Search algorithms have to handle repeated states to be efficient:

- A **tree search algorithm** (naively implemented) may get stuck in infinite loops or do redundant work because it doesn't recognize that it's revisiting a state. For instance, DFS on a naive tree representation of the 8-puzzle might keep sliding a tile left and right forever.
- A **graph search algorithm** uses a memory of visited states (often called a *closed list* or explored set) to avoid revisiting them. This prevents loops and trims the search space by not exploring the same configuration twice. Both the DFS and BFS pseudocode in our notes include steps to eliminate or skip children that have been seen before (in either open or closed lists) – that's the graph search aspect ensuring each state is processed once.

In an exam context, you should mention that **using a graph representation with visited-state checking is crucial for correctness and efficiency** in many problems. It's a commonly overlooked detail: many students implement search as if on a tree and forget to handle repeated states. Remember:

- In a **tree search**, the branching factor and depth determine the number of nodes, and it can blow up enormously because it may count duplicates multiple times.

- In a **graph search**, the effective branching reduces as you prune revisits, but you need extra memory to track visited nodes.

**State Representation Choices:** The way you represent a state can affect the branching factor and efficiency. For example, in the 8-puzzle, representing moves as “move blank up/down/left/right” (instead of moving individual numbered tiles) means each state has up to 4 successors (the blank’s possible moves) and avoids treating swapping two tiles as separate actions. This clever representation (moving the blank) keeps the branching factor lower (around 2-3 on average) compared to if we allowed any tile to move into the blank (which would double-count moves). The notes highlight that *“the greater the number of operators, the larger the branching factor”* and thus **operators should be chosen carefully** to minimize redundant moves.

Other considerations in state space representation include:

- **What to encode in a state description?** Only include what’s necessary for decision-making. (E.g., in TSP a state might encode the set of cities visited so far and the current city, rather than the exact path taken, to avoid state explosion.)
- **Classification of problem types:** Are we looking for *any* solution (just need to reach a goal state), or *optimal* solution (best path cost)? Some problems only need a feasible solution (like n-Queens or a logic puzzle), whereas others like TSP demand the best possible solution. This affects whether we use algorithms that guarantee optimality or if a “good enough” approach suffices.
- **Solution path vs final state:** Some problems (like a puzzle) require the path (sequence of moves) as the answer, not just the final configuration, so the search must record or output the moves taken. Others just require a final state that satisfies criteria (like a solved board configuration, or a proof of existence).

In summary, properly representing the problem as a state space is the first step in applying search. It involves deciding on how to encode states, what the moves are, and recognizing if the state space is **tree-structured or graph-structured**. A well-posed state space and understanding of its structure will guide which search algorithm is most suitable and how to implement it without running into unnecessary inefficiencies.

*(A practical insight: Often in AI, the hardest part is \_formalizing\_ a messy real-world problem into a clean state space model that algorithms can then solve. Once you have a good representation, the rest becomes an algorithmic choice.)\_*

# Uninformed Search Strategies (Blind Search)

**Uninformed search** algorithms explore the state space without any domain-specific guidance; they treat all states as equally promising if not yet explored. These strategies guarantee finding a solution if one exists (provided the state space is finite or certain precautions are taken) but can be very inefficient because they do not have “insight” into which paths might be more fruitful. The primary uninformed methods you should know are **Depth-First Search**, **Breadth-First Search**, and **Iterative Deepening**:

- **Depth-First Search (DFS):** This strategy explores a path as deeply as possible before backtracking. In other words, it expands the newest (most recently discovered) unexpanded state first, diving into one branch of the tree/graph. Implementationally, DFS uses a stack (LIFO order) for the frontier (the “open” list of discovered states).

*Characteristics:* DFS is **memory-efficient** (it only needs to store a path from the root to a leaf plus unexplored siblings – memory grows linearly with depth,  $O(d)$  for depth  $d$ , ignoring visited-set which if used makes it  $O(n)$  with  $n$  states). It can find a solution without examining all shallower states if the solution happens to lie deep in one branch. However, DFS has some downsides: it is **not guaranteed to find the shortest path** (it may find a long solution in one deep branch even if a shorter solution exists elsewhere). Also, a naive DFS can **fail to terminate** if it gets caught in an infinite branch or cycle (hence incorporating visited state checks or depth limits is important).

*Use cases:* DFS is often fine for problems where solutions are known to be deep or when we just need *some* solution regardless of optimality. For example, solving the Towers of Hanoi with DFS works because the solution is at a maximum depth and we don't mind the length (we just want the moves). In a huge space, DFS might be the only feasible way to go arbitrarily deep (since BFS would run out of memory).

*Theoretical notes:* Without a depth bound, DFS is *incomplete* in infinite search spaces. With cycle checking (graph search) it's complete for finite spaces. Time complexity in worst case is  $O(b^m)$  where  $b$  is branching factor and  $m$  is maximum depth (which could be exponential). Typically, exam answers mention that DFS can be very fast if it gets lucky but also that it can miss shallow solutions or loop indefinitely if not careful.

**Insight (Analogy):** DFS is like exploring a maze by always turning left and going as far as you can until hitting a dead end, then backtracking – you might find an exit deep in, but you could also wander a long corridor unnecessarily while an exit was just around the first right turn.

- **Breadth-First Search (BFS):** BFS explores the state space level by level. It expands the oldest nodes in the frontier first (FIFO order, using a queue). That means it first explores all states one move away from the start, then all states two moves away, and so on.

*Characteristics:* BFS is **guaranteed to find the shortest solution** (in terms of number of moves) if all actions have equal cost, because it explores by increasing path length. It is **complete** for finite state spaces (it will eventually exhaust all possibilities at depth  $d$

before moving to  $d+1$ ). The major drawback of BFS is its **high memory usage** – it stores every frontier node at a given level (which can be exponential in depth). For a branching factor  $b$  and depth  $d$  solution, the memory (and time) complexity is about  $O(bd)O(bd)$ . This can quickly become infeasible as  $d$  grows. For example, if each state has  $b=3$  successors, at depth 10 BFS may have to handle  $3^{10} \approx 59k$  nodes; at depth 15,  $3^{15} \approx 14\text{million}$  nodes, which is a lot to store in memory. BFS can also be slow if the solution is deep, because it does not prioritize likely paths – it exhaustively combs all shallow states even if they are irrelevant.

*Use cases:* BFS is ideal when the search depth is small or when we absolutely need the **shortest path** solution (e.g., finding the smallest number of moves in a puzzle, or the shortest route on an unweighted map). It's also conceptually simple and guarantees optimality in those conditions.

*Theoretical notes:* BFS's completeness and optimality (for equal costs) are its key selling points. On an exam, a common point is that BFS might **blow up (combinatorial explosion)** due to the branching factor – a phrase to remember.

**Real-world example:** BFS is analogous to exploring all locations at distance 1, then distance 2, etc. If you were trying to find a friend in a building floor by floor, BFS is like checking every room on the first floor, then every room on the second, etc., ensuring you find them in the least number of rooms opened, but possibly opening a lot of unnecessary rooms on the way.

- **Depth-First Iterative Deepening (DFID or IDDFS):** This strategy combines the benefits of DFS and BFS. The idea is to run a DFS with a depth limit of 0 (i.e., just the start), then restart DFS with depth limit 1, then 2, and so on, increasing the limit until the goal is found. Essentially, it performs DFS in layers, mimicking BFS in terms of finding the shallowest solution, but using DFS's space-saving nature for each layer.

*Characteristics:* IDDFS is **complete** and will find an optimal solution (optimal in terms of shortest path length) if one exists, assuming uniform step costs, just like BFS. Its time complexity is similar to BFS in the worst case  $O(bd)O(bd)$ , but with some overhead due to re-exploring the same nodes at each new depth bound. However, this overhead is not as bad as it sounds: for example, exploring depth 0 + depth 1 + ... + depth  $d$  results in exploring the shallow nodes multiple times, but the cost is geometric series like  $b^0 + b^1 + \dots + b^d \approx O(b^d)$  for large  $d$ . If  $b$  is not huge, the re-searching cost is acceptable. The **memory usage** of IDDFS is much lower than BFS – it only needs memory proportional to the depth (like DFS) for the current iteration.

*Why use IDDFS:* It's extremely useful when the search space is very large or infinite-depth, because plain BFS would run out of memory, and plain DFS might miss the shallow solution. IDDFS gives you the depth-guaranteed approach of BFS with the space of DFS. It's often the default for finding shortest paths in an unknown state space when you don't have heuristics.

*Theoretical notes:* IDDFS returns the same solution as BFS would (the shallowest one) but

with a modest time overhead. It's complete in finite spaces or infinite spaces with finite branching and at least one solution at finite depth. On an exam, you might be asked "Why use iterative deepening instead of BFS or DFS?" – the answer: it achieves BFS's guarantee of finding the shortest solution while drastically reducing memory use, making it feasible to search deeper. The trade-off is performing repeated work, but typically that cost is outweighed by the memory savings.

- (*Additional uninformed strategies*: Other uninformed strategies include **Uniform-Cost Search** (UCS), which is like BFS but for weighted graphs – always expanding the least-cost node first. It guarantees optimality for general step costs. UCS wasn't explicitly listed in the question, but just know it's essentially Dijkstra's algorithm for shortest paths and can be seen as an uninformed search that uses path cost. Also, **Bidirectional Search** where you search from both start and goal towards the middle can drastically cut down search depth if you have a goal state explicitly known. These might be beyond Chapter 3 scope if not mentioned, but they are good to know for completeness.)\*

**Comparing DFS and BFS:** A common exam question is to compare these two:

- *Completeness*: BFS is complete (in finite space), DFS is not guaranteed (unless modified).
- *Optimality*: BFS is optimal for unweighted shortest paths; DFS is not (it may find a non-optimal solution first).
- *Time Complexity*: Both can be exponential in worst case, but BFS explores all shallower nodes first, DFS can sometimes find a solution faster if it's deep. BFS examines all states up to depth  $d$  for a depth- $d$  solution; DFS could get lucky or unlucky.
- *Space Complexity*: BFS requires storing all nodes at the current frontier (which is  $O(b^d)$  in worst case). DFS stores one path and maybe some neighbors ( $O(b \cdot d)$  at most for a path depth  $d$  in tree search). This is the key reason to prefer DFS or IDDFS when memory is tight.
- *When to use*: If you know the solution isn't too deep and need the shortest path -> BFS. If the space is huge or infinite and you just need a solution or need to go deep -> DFS (with care) or IDDFS for safety.

Understanding these pros/cons and being able to articulate them clearly (possibly as bullet points or a small table) is very useful for exams. Most importantly, tie them back to consequences: e.g., "BFS might run out of memory first, DFS might run out of time (in a bad case) or fail to find a solution if not careful."

## Informed (Heuristic) Search Strategies

Uninformed methods treat all nodes blindly, but **informed search** uses extra knowledge in the form of a **heuristic function** to estimate how close a given state is to a goal. The heuristic guides the search to explore more promising paths first. This often drastically improves

efficiency, at the cost of potentially sacrificing completeness or optimality if the heuristic is not perfect. Key heuristic search strategies to know include **Greedy Best-First Search**, **Hill-Climbing**, **A\***, and **Branch and Bound** techniques:

- **Greedy Best-First Search:** This algorithm expands the node that appears *closest* to the goal, as measured by a heuristic function  $h(n)$ . Typically,  $h(n)$  is an estimate of the cost (or distance) from state  $n$  to a goal. Greedy best-first search selects the next node with the smallest  $h$  value (i.e., most “promising” looking state). For example, in a pathfinding problem on a map,  $h(n)$  might be the straight-line distance from node  $n$  to the destination, and the algorithm will always head in the direction that seems to reduce this distance.

*Characteristics:* Greedy best-first is **fast and can find a solution quickly** because it uses intuition about which direction the goal lies. However, it is **not guaranteed to find an optimal solution** – it might take a path that looks good early but actually is longer or even leads to a dead end. It’s also not guaranteed to be complete in some cases (for instance, if the heuristic is misleading or if there are infinite states that look better but never lead to a goal, it could keep exploring them). In practice, it often finds a solution faster than BFS or DFS, but that solution might be suboptimal.

*Example:* Consider a maze where the goal is on the east side. A greedy strategy using straight-line distance will always try to move eastward. If there’s a wall directly between the start and goal, greedy might keep following along the wall thinking any step toward the goal (eastward) is best, potentially going into a dead-end corner, whereas the optimal path might involve going away (north, then east around the wall). Greedy won’t realize the detour is needed because the heuristic misleads it.

*Use cases:* Greedy best-first is useful when you want a quick and usually decent solution and don’t need optimality. It’s a component in algorithms like hill-climbing and also used in pathfinding when a rough route is acceptable.

*Note:* In terms of implementation, best-first search uses a priority queue (ordered by  $h(n)$ ) for the frontier. It’s like BFS but prioritized by a heuristic instead of depth. On an exam, if asked about “best-first search,” clarify whether it’s the greedy kind (using  $h$  only) or a general concept (sometimes “best-first” can refer to any heuristic-driven search including A\*). Typically, they mean greedy best-first.

- **Hill-Climbing (Local Search):** Hill-climbing is a variant of best-first that is usually described as *iteratively improving a single current state* rather than maintaining a full search frontier. It starts from an initial state and repeatedly moves to a neighboring state that has a better evaluation (higher value or lower cost, depending on formulation) than the current one, until no improvement can be found. In effect, it “climbs the hill” of a quality landscape, always going upward (for maximization problems) or downward (for cost minimization).

*Characteristics:* Hill-climbing is conceptually simple and memory-efficient (it doesn’t maintain a search tree, just the current state and maybe some local info). It’s appropriate for **optimization problems** where the goal is to find a state that maximizes some objective.

However, it suffers from the well-known issue of **local maxima** (or minima) – if it reaches a state that is better than all its neighbors but not the best overall (a local peak), it stops there, thinking it's at the top. This means hill-climbing can get stuck in suboptimal solutions. Other pitfalls include **plateaus** (flat areas where many states have equal value, causing the algorithm to randomly wander) and **ridges** (where the path to a better state isn't directly uphill, so greedy steps can't follow it).

*Variants:* To overcome local maxima, there are variants like **stochastic hill-climbing** (choose randomly among better moves, adding some randomness), **random restarts** (run hill-climbing multiple times from random initial states, to hope one finds the global max), or **simulated annealing** (occasionally allow downhill moves to escape local maxima, analogous to cooling metal annealing – though simulated annealing itself might be beyond our chapters, it's a known variant).

*Analogy:* Hill-climbing is often compared to trying to reach the highest point in a foggy landscape by always moving upward. If you're on a hill that isn't the tallest, you'll climb to that hill's peak and stop, because you can't see a higher point without first going down. This analogy perfectly encapsulates the local maximum problem.

*Use cases:* Hill-climbing is common in tasks like optimizing continuous functions or tuning parameters. It's not guaranteed to find the best solution, but it's fast and simple. For exam answers, note that it doesn't guarantee an optimal solution and isn't systematic, but it's useful for quick approximate solutions. Also, emphasize it's a *local* search method – it doesn't explore alternatives once it commits to a direction (unlike A\* which keeps a whole frontier of options).

*Real-world parallel:* Many human problem-solving scenarios use hill-climbing: for example, iterative improvement strategies like trying to improve a design by small tweaks, or even something like solving a crossword puzzle by filling in best-guesses letter by letter (you might end up with a partially filled puzzle that has no solution – a local max – and need to backtrack or start over).

- *\_A Search: The A algorithm is the gold standard of heuristic search for finding optimal solutions. It extends best-first search by considering both the cost to reach a node and the heuristic estimate to get to the goal. Specifically, A uses an evaluation function  $f(n)=g(n)+h(n)$ , where  $g(n)$  is the cost from the start to node  $n$  (so far) and  $h(n)$  is the heuristic estimate from  $n$  to the goal. By always expanding the node with lowest  $f(n)$  next, A effectively finds the least-cost solution path while using  $h(n)$  to guide the search efficiently.*

*\_Optimality and Admissibility: A is **complete** (it will find a solution if one exists, assuming finite branching and positive costs) and **optimal** – if the heuristic  $h(n)$  is **admissible**. An **admissible heuristic** is one that never overestimates the true cost to reach the goal (it's always optimistic or exact). Under that condition, A will never ignore a truly optimal path in favor of a worse one, because it will eventually explore the optimal path before any path that would exceed the optimal cost. Additionally, if the heuristic is*



also **consistent** (monotonic), *A has the nice property that the first time it pops a goal from the open list, that is the optimal solution and it never needs to revisit states. (Consistency means  $h(n) \leq c(n, n') + h(n')$  for any neighbor  $n'$  of  $n$  with cost  $c(n, n')$ ; essentially, the heuristic estimate obeys the triangle inequality. Consistency implies admissibility and ensures  $f(n)$  for nodes on a path is non-decreasing.)*

*Efficiency: With a good heuristic, A dramatically reduces the number of states explored compared to BFS. In the best case, if  $h(n)$  is very accurate (close to the true cost), A will head almost straight to the goal. In the worst case, if  $h(n)$  is zero (no help at all), A becomes equivalent to uniform-cost search (which is like BFS for equal costs). If  $h(n)$  is too optimistic, A is still optimal but may explore more than necessary;*

*if  $h(n)$  overestimates (inadmissible), A can find suboptimal solutions faster but then it's no longer guaranteed optimal.*

*Memory note: A's big downside is it stores a lot of nodes (everything it's expanded or is waiting to expand), so memory can be a limiting factor. There are variants like IDA (Iterative Deepening A) to mitigate that by using iterative deepening on increasing  $f$ -cost thresholds, but that's more advanced.*

*Example: In pathfinding on a grid, if we use Manhattan distance as  $h$ , A will essentially perform a guided BFS toward the goal. For the 8-puzzle, classic heuristics like "number of misplaced tiles" or "Manhattan distance of all tiles" are admissible. If the puzzle can be solved in 20 moves, a brute force might explore millions of nodes, but A with Manhattan distance might explore only thousands, because it focuses on moves that actually reduce the distance of tiles to where they belong.*

*Usage: A is used in countless applications: route planning (GPS finds shortest routes using A variations), puzzle solvers, game AI (to navigate characters through a map), network routing, and more. It's popular because it finds the optimal solution and is often the most efficient way to do so.*

*Exam pointers: Be ready to define what  $g, h, fg, h, f$  are. Emphasize why admissibility matters. Possibly mention that if  $h=0$  then A is just uniform-cost search, and if  $h$  is admissible A will never find a costlier solution before a cheaper one. You might also mention complexity: in worst-case scenarios, A still runs in exponential time (because there might be an exponential number of states to examine if the heuristic isn't sufficiently pruning), but in practice a good heuristic can make it feasible. Another angle: A is like a **balance** between BFS and greedy – BFS cares only about  $g$  (distance from start), greedy cares only about  $h$  (distance to goal), and A smartly combines them. This balance is why  $A^*$  is generally superior to either alone.*

- **Branch and Bound:** Branch-and-bound is a general **search optimization technique** rather than a specific algorithm, but in the context of AI search problems it refers to exploring branches of the search tree while keeping track of the best solution found so far (the current bound) and **pruning** any branch that cannot possibly yield a better solution than the best one currently known. In practice, it's often used with depth-first search for

optimization problems. For example, in a DFS through a state space (like for TSP or other optimization problems), if you accumulate a path cost and it exceeds the best solution's cost found so far, you can stop exploring that path (prune it) because it's not going to be optimal. *How it works:* Each partial solution has a *bound* (which could be the partial cost so far plus a heuristic estimate of completion). If this bound is worse than an existing complete solution's cost, you cut off that branch. This requires being able to estimate a lower bound on any completion of the current partial path (so you know if it can't possibly beat the best). In TSP, for example, a simple bound for a partial tour is cost so far plus a straight-line heuristic to complete, or more sophisticated, the cost so far plus minimum spanning tree on remaining cities, etc.

*Relation to A:* If you think about it, A is a kind of best-first branch-and-bound: it always expands the lowest  $f=g+hf=g+h$  (that's like a lower bound on complete path cost) first, so any path with  $f$  higher than the best solution found will naturally get pruned (or rather, never expanded until it possibly becomes best). Branch-and-bound in a DFS context might expand one complete solution (find some solution cost  $C$ ), then use  $C$  as the bound to prune other branches. It's often combined with DFS because DFS goes deep to find a solution quickly which gives you an initial bound, then continues to explore to see if something better can be found, pruning along the way.

*Use cases:* Branch-and-bound is widely used in solving combinatorial optimization problems like TSP, knapsack, job scheduling, etc., in more advanced algorithm courses. In AI search context, it underscores the importance of **pruning** – not exploring parts of the space that you know won't lead to a better answer.

*Exam focus:* Understand that branch-and-bound is about **keeping track of the best found solution and using it to cut off inferior paths**. It guarantees that by the end, the best found is optimal because anything better would have had to survive the bound check, which none did. It's complete and optimal, but like other searches can still be exponential in worst case (if few branches get pruned). A detail students sometimes miss: branch-and-bound requires a way to compare partial solutions to a known complete solution – meaning you need some optimistic estimate for any partial path (which is essentially a heuristic lower bound). Without a good bound function, it doesn't prune much.

*In summary:* Branch-and-bound improves efficiency by **avoiding exploration of costly paths**, analogous to a clever chess player pruning moves that already look worse than a scenario she can achieve via another line of play.

In practice, heuristic search algorithms aim to reduce the **combinatorial explosion** of brute-force search by leveraging domain knowledge through heuristics. The trade-off is that if heuristics are poorly chosen, they can mislead search or provide little benefit. But when well-chosen, they are extremely powerful. For example, using heuristics is the only reason we can solve large 15-puzzles or complex routing problems that would be infeasible with pure BFS.

*(Human cognition insight: Humans often use heuristics in problem-solving – we don't try every possibility when solving a puzzle, we use intuition about what "looks" closer to the solution. This is exactly what heuristic search automates. However, just like a person's intuition can be wrong or get stuck (like a misleading puzzle move that seems good but isn't), an AI's heuristic can lead it astray. Balancing exploration and informed guidance is key, and algorithms like A\* manage this balance optimally under the right conditions.)\**

## Heuristics: Construction and Admissibility

Heuristics are at the heart of informed search, and constructing a good heuristic is as much an art as it is a science. Let's break down what you need to know about heuristics, especially focusing on the concept of **admissibility** (since exam questions love to ask about that):

### What is a Heuristic?

A **heuristic** is a function  $h(n)$  that estimates the cost (or distance, or steps) from a state  $n$  to the nearest goal state. It provides extra knowledge beyond the problem definition to help guide search. A heuristic is *problem-specific*. For example:

- In the 8-puzzle, a common heuristic is the **Manhattan distance**: for each tile, compute how far (in grid moves) it is from its goal position, and sum these distances. This gives a number that *estimates* how many moves remain to solve the puzzle from the current configuration.
- In a navigation problem (like Google Maps routing), a heuristic from any location could be the straight-line (Euclidean) distance to the destination – it's an estimate of remaining driving distance (assuming you could go in a straight line).
- For the Traveling Salesman Problem, a heuristic could be something like the cost of a Minimum Spanning Tree connecting all unvisited cities from the current city – that's a plausible underestimate of the remaining tour length.

### Admissible Heuristics:

A heuristic  $h(n)$  is called **admissible** if it never overestimates the true minimal cost from  $n$  to a goal. In other

words,  $h(n) \leq (\text{actual lowest cost from } n \text{ to goal})$  for all states  $n$ . Admissibility is crucial for algorithms like A\* to *guarantee an optimal solution*.

*Intuitively, an admissible heuristic is optimistic – it might think the goal is closer than it actually is, but it's never overly optimistic (never guesses lower cost than actually possible). Because it underestimates or exactly estimates the remaining cost, A\* will never ignore a path thinking it's too expensive when in reality it was the cheapest – the true cheapest path will always look at least as promising as it really is, so A\* will pursue it.*

*Example of admissible vs inadmissible:* In the 8-puzzle:

- **Misplaced Tiles count:** The number of tiles not in their goal position. This is admissible because any tile out of place needs at least one move to put in place, so total moves needed is at least that number.
- **Manhattan Distance:** Admissible because every tile must move at least its Manhattan distance to reach home (it can't be done in fewer moves than that distance).
- An **inadmissible** heuristic would be something like Manhattan distance *1.5*. *That might often overshoot the true moves required, for instance estimating 15 moves when perhaps it can be solved in 12. If you feed that into A*, it might prune or overlook the optimal path due to thinking it's costlier than it is, and thereby find a suboptimal solution faster.

Why might someone use an inadmissible heuristic? Because sometimes a slightly greedy overestimating heuristic will expand far fewer nodes (making the search faster), albeit at risk of missing the optimal answer. In practice, if we care about optimality, we stick to admissible (or at least hope the overestimation is small enough not to skip the optimal path).

### **Consistent Heuristics:**

Another property often mentioned is **consistency** (or monotonicity). A heuristic is consistent if for every state  $n$  and every action from  $n$  to a successor  $n'$  with step cost  $c(n, n')$ , it holds that  $h(n) \leq c(n, n') + h(n')$ . This basically means the estimated cost from  $n$  is no more than the step cost plus estimated cost from  $n'$ . Consistency implies that as you move from a state to any neighbor, your heuristic estimate doesn't drop by more than the actual cost of that step. This property ensures that  $f(n) = g(n) + h(n)$  for  $A$  is *non-decreasing along any path* (it's like a triangle inequality for costs). *If heuristics are consistent, A* will never need to reopen a node (i.e., once it's visited a state with the lowest cost, that's the best way to get there). Most admissible heuristics used in practice are also consistent (all the common puzzle heuristics like Manhattan distance are consistent). For exam purposes, if not explicitly asked, you can remember consistency as a sufficient condition for easier implementation, but admissibility is the main condition for optimality.

### **Designing Heuristics – Common Techniques:**

- **Relaxed Problems:** A common way to derive an admissible heuristic is to consider a simplified (or relaxed) version of the problem that is easier to solve. The cost to solve the relaxed problem from state  $n$  is an admissible heuristic for the original problem, because the relaxed problem is, by definition, no harder than the original – so its solution cost is a lower bound. For example, for the 8-puzzle, a relaxation is “tiles can move anywhere freely”

– solving that is just placing each tile correctly which leads to the misplaced tile count heuristic. Another relaxation: “tiles can move without regard for others, as if they could pass through each other” – solving that gives the Manhattan distance (each tile moves straight to its spot).

- **Heuristic Combination:** If you have multiple admissible heuristics  $h_1, h_2$ , their maximum  $h_{\max}(n) = \max(h_1(n), h_2(n))$  is also admissible (it's never an overestimate if each isn't, and it's closer to the true cost than either alone typically). Sometimes combining heuristics this way gives a stronger estimate.
- **Pattern Databases:** For certain puzzles, people precompute the exact cost for a subset of the puzzle (a subproblem) for all configurations, and use that as a heuristic. This is an advanced idea but notable: for instance, in the 15-puzzle, you can precompute the minimum moves needed for each configuration of, say, 7 tiles (ignoring the others), store these in a database, and use that sum as a heuristic. Those can be very accurate and still admissible.
- **Domain Knowledge:** Use specific insights. For TSP, an admissible heuristic is perhaps “take the minimum edge out of the current city plus minimum spanning tree of remaining cities” – it's a bit complex but it's rooted in domain knowledge that a spanning tree is the shortest way to connect remaining points (no cycles).
- **Learned Heuristics:** In modern AI, sometimes heuristics are learned by machine learning on similar problem instances. But that's beyond our scope; just interesting that neural nets can even learn to approximate heuristic values, for example.

### Heuristics in Practice:

A well-known result: *more informed* (higher valued, yet admissible) heuristics lead *A* to expand fewer nodes. However, there's a point of diminishing returns – computing a very accurate heuristic might itself be costly. A trivial heuristic is  $h(n)=0$  (which is admissible but gives no guidance). A perfect heuristic  $h^*(n)$  that gives the true cost to goal would make *A* expand only the optimal path nodes – best case – but computing that is as hard as solving the problem, so it's not feasible. We seek a middle ground: easy to compute but informative.

For exam preparation, focus on:

- **Definition of admissible heuristic** – likely a one-liner like “a heuristic that never overestimates the true cost to reach the goal.”
- **Significance in  $A^*$**  – that it guarantees  $A^*$  finds optimal solutions.
- Possibly give an example of constructing a heuristic or mention one. E.g., “For the 8-puzzle, one admissible heuristic is the sum of Manhattan distances of each tile from its goal position.”
- Perhaps mention that heuristics are **non-negative** typically and  $h(\text{goal})=0$  (a requirement for admissibility and logical correctness).

- If asked about creating a heuristic: explain the relaxed problem technique or give an intuitive design.

**Common Mistake to Avoid:** Some students think any intuitive guess is fine; but if it overestimates,  $A^*$  can fail. For instance, a student might think in an 8-puzzle, “maybe multiply Manhattan distance by a factor to speed it up” – that breaks admissibility. Or forgetting the requirement  $h(\text{goal})=0$  (goal state should have zero estimated cost to finish, since you’re already there). These details might appear as trick questions.

**Real-world tie-in:** Heuristics pervade human decision-making; for example, a heuristic for solving a Rubik’s Cube might be “first get all corner pieces in place” as a subgoal – it’s not a strict rule, but it guides the approach. Admissibility in human terms means you never think a task is easier than it is – easier to err on thinking it’s a bit easier (optimistic) than underestimating (which could cause failure to plan properly).

By mastering how heuristics are formed and used, you can appreciate why certain AI algorithms are so powerful. The combination of algorithm (like  $A^*$ ) and a well-chosen heuristic can crack problems that are otherwise intractable by brute force.

## Exam Wildcards (Frequently Overlooked Details)

- **Turing Test Setup:** Remember that the original Turing Test (Imitation Game) involves *two humans and one machine*. A common mistake is to think it’s just one human vs one machine. Actually, one human judge interrogates both a hidden human and a hidden machine; the goal is for the machine to be indistinguishable from the human. This detail is often overlooked but is key to describing the test correctly.
- **Chinese Room’s Point:** The Chinese Room isn’t claiming the man inside “understands Chinese” – in fact, it’s the opposite. It’s illustrating that *syntax (rule following)  $\neq$  semantics (understanding)*. Some students misinterpret it as “the room as a whole understands Chinese” (Searle would argue it doesn’t). The wild insight here: even if a computer passes the Turing Test, it might still lack true understanding, just as the person in the Chinese Room manipulates symbols without comprehension.
- **Difference between AI Search Trees vs. Problem Graphs\*\*\*:** *In search algorithms, especially when writing pseudocode, many forget to handle repeated states. The distinction between a tree search (which might revisit states endlessly) and a graph search (which tracks visited states) is crucial. Always mention the need for a closed list or visited set in algorithms like DFS/BFS/A to avoid infinite loops and redundant work.*

- **DFS Space Complexity Advantage:** It's easy to recall that BFS finds the shortest path, but don't forget that **DFS's big advantage is space efficiency**. On an exam, when comparing BFS and DFS, many mention time or completeness but omit that BFS might exhaust memory on large breadth, while DFS keeps memory linear in depth. This is often a point instructors look for in answers.
- **Iterative Deepening == Optimal (for unit costs):** Iterative Deepening DFS will find the same solution as BFS (the shallowest solution). Some answers mistakenly think IDDFS might give a longer solution due to DFS nature – in reality, it systematically increases depth and will find the shortest solution first. Also, note that the re-exploration overhead is not as terrible as it seems (exploring depth 1 + depth 2 + ... depth d is about 11% overhead compared to BFS, if branching factor is around 2; even with higher branching, overhead stays reasonable).
- **Heuristic = 0 case:** An often overlooked point is that algorithms like A and greedy best-first generalize other algorithms. *If your heuristic  $h(n)$  is zero for all nodes, A essentially becomes uniform-cost search (or BFS if all step costs equal).* This underscores that heuristics only help A; *in the worst case it's as bad as BFS, never worse. It's a good sanity check: A can only improve (or tie) performance vs uninformed search, never degrade it, if implemented properly.*
- **Admissible vs Consistent:** Many students confuse these or think they are interchangeable. Remember: *admissible* means never overestimates (focuses on goal accuracy), *consistent* means heuristic values obey the triangle inequality (focuses on local transitions). Consistency implies admissibility. If a heuristic is inconsistent, A\* still works but might have to reopen nodes, which complicates things. On an exam, if asked, give a brief distinction and perhaps note that most common heuristics (like for puzzles) are consistent.
- **Expert Systems vs. Learning Systems:** A subtle point is the difference between how expert systems acquire knowledge versus systems like neural networks. Expert systems rely on **knowledge engineering** – humans write the rules. Neural networks and other learning systems rely on data. This difference often leads to wildcards in exam questions, like asking for pros/cons of expert systems (e.g., they can justify conclusions via explicit rules but struggle if rules are incomplete, whereas neural nets can generalize from examples but act as a “black box” with no clear justification).
- **Classification of Problems:** Recall the note that some problems require just a solution, others need *the best* solution. For instance, n-Queens or a logic puzzle just needs one valid arrangement (any solution will do), whereas an optimization like TSP needs the optimal route. This affects what algorithm you choose (e.g., you might use DFS or heuristic search to just find any solution quickly, but for an optimal solution you'd use A\* or branch-and-bound). It's easy to overlook that distinction, but if a question asks “what search algorithm would you use and why,” make sure to tie it to whether we need optimality or just a feasible solution.

- **Real-world Context in Answers:** Students often focus on mechanics and forget examples. Including a quick real-world tie-in (like mentioning GPS for A\* or a medical expert system example) can be a wildcard that makes an answer stand out. It shows you understand not just the abstract concept but also its application. Just ensure the example is accurate and relevant.